


one query, every machine

01 Ask once – every node answers



02 One cluster, not many



We run one cluster. The other way is many clusters, joined together by cross-cluster search.

PROS OF MANY CLUSTERS

- + One bad query cannot slow the other clusters.
- + Each cluster upgrades on its own schedule.

PROS OF ONE CLUSTER

- + Built-in cross-cluster search does not scale to 100 machines.
- + The cluster places shards and replicas itself, across all machines.
- + One set of users, roles and saved searches.
- + Detection and alerting streamlined.
- + No federation layer of our own to write and keep alive.

04 Fanout – not everything is perfect

Every machine pays

One search asks all of them, so all of them spend processor time on it. It is as slow as the slowest machine.

Reconstruction comes first

Those same machines reconstruct events. That is their real job, and log search must not compete with it.

So we ask only when we must

The live lane covers watching. But “which machines saw this error in run 561234?” has to ask all of them.